

Kubernetes 101

Tommy Chen @ TechCCU 2016



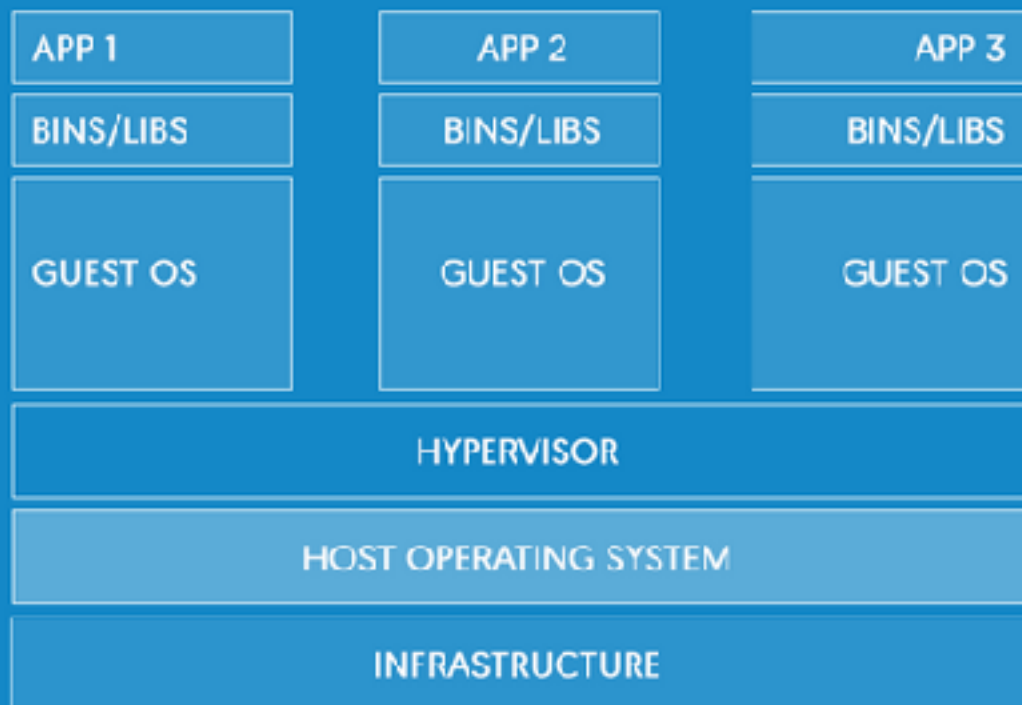
Tommy Chen

- a.k.a. SkyArrow
- Dcard / Full Stack Developer
- @tommy351



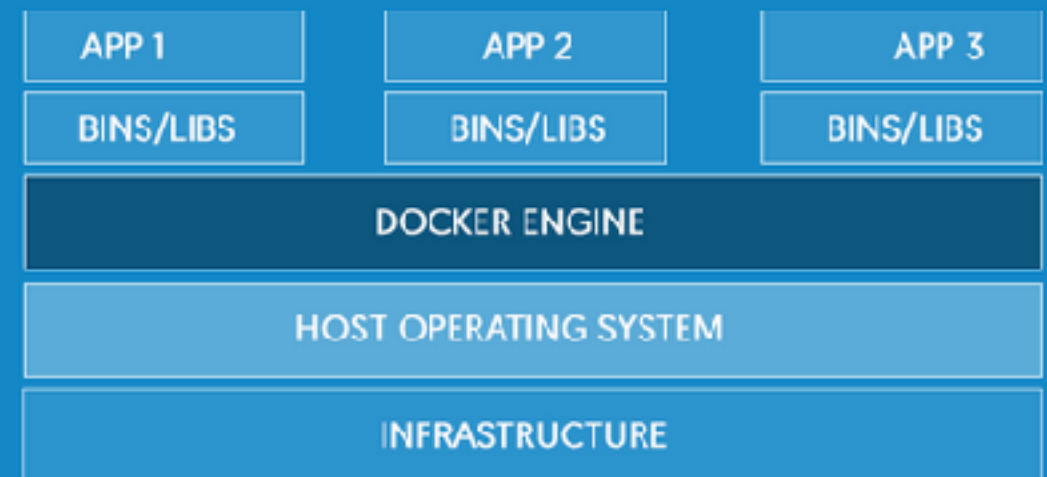
1

Recap Docker



VIRTUAL MACHINES

Virtual machines include the application, the necessary binaries and libraries, and an entire guest operating system -- all of which can amount to tens of GBs.



CONTAINERS

Containers include the application and all of its dependencies --but share the kernel with other containers, running as isolated processes in user space on the host operating system. Docker containers are not tied to any specific infrastructure: they run on any computer, on any infrastructure, and in any cloud.

<https://www.docker.com/what-docker>

Dockerfile

```
FROM scratch  
COPY app /  
EXPOSE 4000  
CMD ["/app"]
```

Build

```
$ docker build --tag tommy351/my-server .
```

```
Sending build context to Docker daemon 3.943 MB
```

```
Step 1 : FROM scratch
```

```
--->
```

```
Step 2 : COPY app /
```

```
---> 4c25e1b50576
```

```
Removing intermediate container 22b9992b6f64
```

```
Step 3 : EXPOSE 4000
```

```
---> Running in 975456111f69
```

```
---> 03095e7edf87
```

```
Removing intermediate container 975456111f69
```

```
Step 4 : CMD /app
```

```
---> Running in 9bc9777f565e
```

```
---> f41f9b6ad7d0
```

```
Removing intermediate container 9bc9777f565e
```

```
Successfully built f41f9b6ad7d0
```

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED
tommy351/my-server	latest	f41f9b6ad7d0	6 seconds ago
3.939 MB			

Run

```
$ docker run -p 4000:4000 -d tommy351/my-server
```

```
9eb06be69a532974e811aa89fdf44805fe63616e684c5c38e7fc  
7d766b43dc50
```

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND
CREATED	STATUS	PORTS
NAMES		
9eb06be69a53	tommy351/my-server	"/app"
21 seconds ago	Up 20 seconds	
0.0.0.0:4000->4000/tcp	fervent_heyrovsky	

```
$ curl http://localhost:4000
```

```
Hello world
```

Push & Pull

```
$ docker push tommy351/my-server
```

```
The push refers to a repository [docker.io/tommy351/my-server]
```

```
8413560c067f: Pushed
```

```
latest: digest:
```

```
sha256:becf447f4dc2a950a10330a86d42209ed42c3a4f7d3636fd50cdd5eeaf358952 size: 528
```

```
$ docker pull tommy351/my-server
```

```
Using default tag: latest
```

```
latest: Pulling from tommy351/my-server
```

```
432847f70861: Already exists
```

```
Digest:
```

```
sha256:becf447f4dc2a950a10330a86d42209ed42c3a4f7d3636fd50cdd5eeaf358952
```

```
Status: Downloaded newer image for tommy351/my-server:latest
```

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED
tommy351/my-server	latest	f41f9b6ad7d0	2 minutes ago
	3.939 MB		

2

What is
Kubernetes?

Why Kubernetes?

- Docker 只能用在單機
- 必須要有軟體來處理多機之間的關係
 - Docker Swarm
 - Kubernetes
 - Marathon (Mesos + DCOS)
 - AWS ECS

功能

- 處理多機之間的關係
- 健康檢查
- 自動水平擴展
- 負載平衡
- 漸進更新 (Rolling update)
- 服務探索
- 設定管理

3

Get
Ready

Minikube

在本機架一個單節點的 Kubernetes cluster，方便進行測試。

<https://github.com/kubernetes/minikube>

```
$ curl -Lo minikube https://storage.googleapis.com/minikube/releases/v0.10.0/minikube-darwin-amd64
```

```
$ chmod +x minikube
```

```
$ sudo mv minikube /usr/local/bin/
```

kubectl

安裝 Kubernetes 的管理工具。

透過 gcloud 安裝：

<https://cloud.google.com/sdk/downloads>

```
$ curl https://sdk.cloud.google.com | bash
```

```
$ exec -l $SHELL
```

```
$ gcloud components install kubectl
```

從 binary 安裝：

<https://coreos.com/kubernetes/docs/latest/configure-kubectl.html>

```
$ curl -O https://storage.googleapis.com/kubernetes-release/release/v1.3.6/bin/darwin/amd64/kubectl
```

```
$ chmod +x kubectl
```

```
$ mv kubectl /usr/local/bin/kubectl
```

啟動 minikube

```
$ minikube start
```

```
Starting local Kubernetes cluster...
```

```
Kubectl is now configured to use the cluster.
```

Hello World

```
$ kubectl run my-server --image tommy351/my-server --port 4000
```

```
deployment "my-server" created
```

```
$ kubectl get deployment
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
my-server	1	1	1	1	17s

```
$ kubectl get replicaset
```

NAME	DESIRED	CURRENT	AGE
my-server-1549122061	1	1	20s

```
$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
my-server-1549122061-9io6v	1/1	Running	0	23s

```
$ kubectl expose deployment my-server --type NodePort
```

```
service "my-server" exposed
```

```
$ kubectl get service
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	10.0.0.1	<none>	443/TCP	22m
my-server	10.0.0.133	<nodes>	4000/TCP	5s

```
$ curl $(minikube service my-server --url)
```

```
Hello world
```


剛剛做了什麼？

- `kubectl run`
 - 建立 Deployment → ReplicaSet → Pod
- `kubectl expose`
 - 建立 Service
 - 暴露服務給外界存取

4

Basic Concepts

常用指令

建立資源

```
$ kubectl create -f <filename>
```

顯示資源列表，type 可以是 pod, rc, service, etc.

```
$ kubectl get <type>
```

取得單一資源的資訊

```
$ kubectl get <type> <name>
```

刪除單一資源

```
$ kubectl delete <type> <name>
```

```
$ kubectl delete -f <filename>
```

在 pod 裡執行指令 (相當於 docker exec)

```
$ kubectl exec -it <pod> <command>
```

常用指令

顯示 pod 紀錄 (相當於 docker logs)

```
$ kubectl logs <pod>
```

更新資源

```
$ kubectl apply -f <filename>
```

顯示 kubernetes 的所有事件

```
$ kubectl get events
```

開啟 Kubernetes proxy，預設在 localhost:8001，

localhost:8001/ui 是 dashboard

```
$ kubectl proxy
```

連接埠轉發 (e.g. kubectl port-forward postgres 5433:5432)

```
$ kubectl port-forward <pod> <local port>:<remote port>
```

Pod

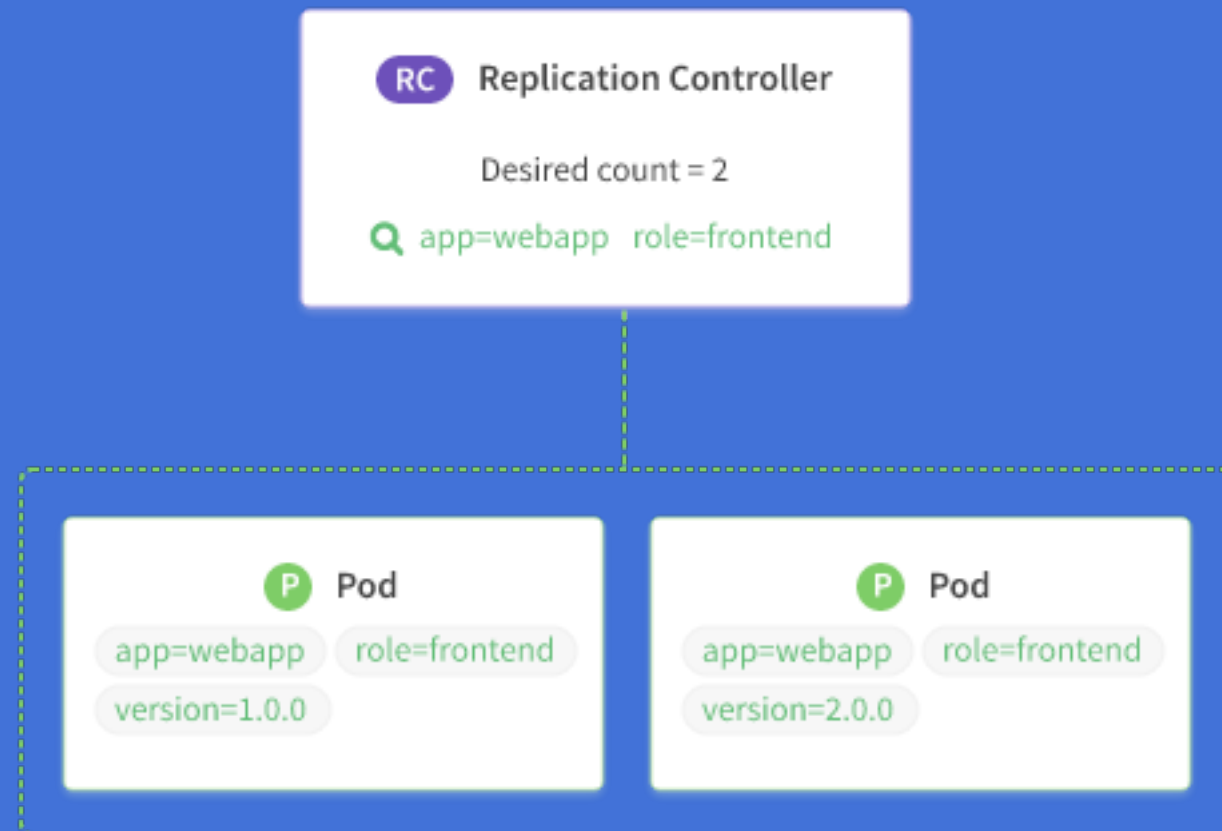
- 一個以上容器的集合
- Kubernetes 中基本的執行單位
- 每個 Pod 都會配一個 Cluster IP
- 死掉的話不會重啟，必須透過 Replication controller 或 Replica set 管理

Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: my-server
  labels:
    name: my-server
spec:
  containers:
    - name: my-server
      image: tommy351/my-server
      ports:
        - containerPort: 4000
```

Replication Controller

- 管理 Pod 的生命週期
- 確保指定數量的 Pod 在執行



Replication Controller

apiVersion: v1

kind: ReplicationController

metadata:

name: my-server

spec:

replicas: 2 Pod 數量

selector:

name: my-server

Selector 會選中 Label 符合的 Pod

template:

metadata:

name: my-server

labels:

name: my-server

和剛剛的 Pod 內容一樣

spec:

containers:

– **name:** my-server

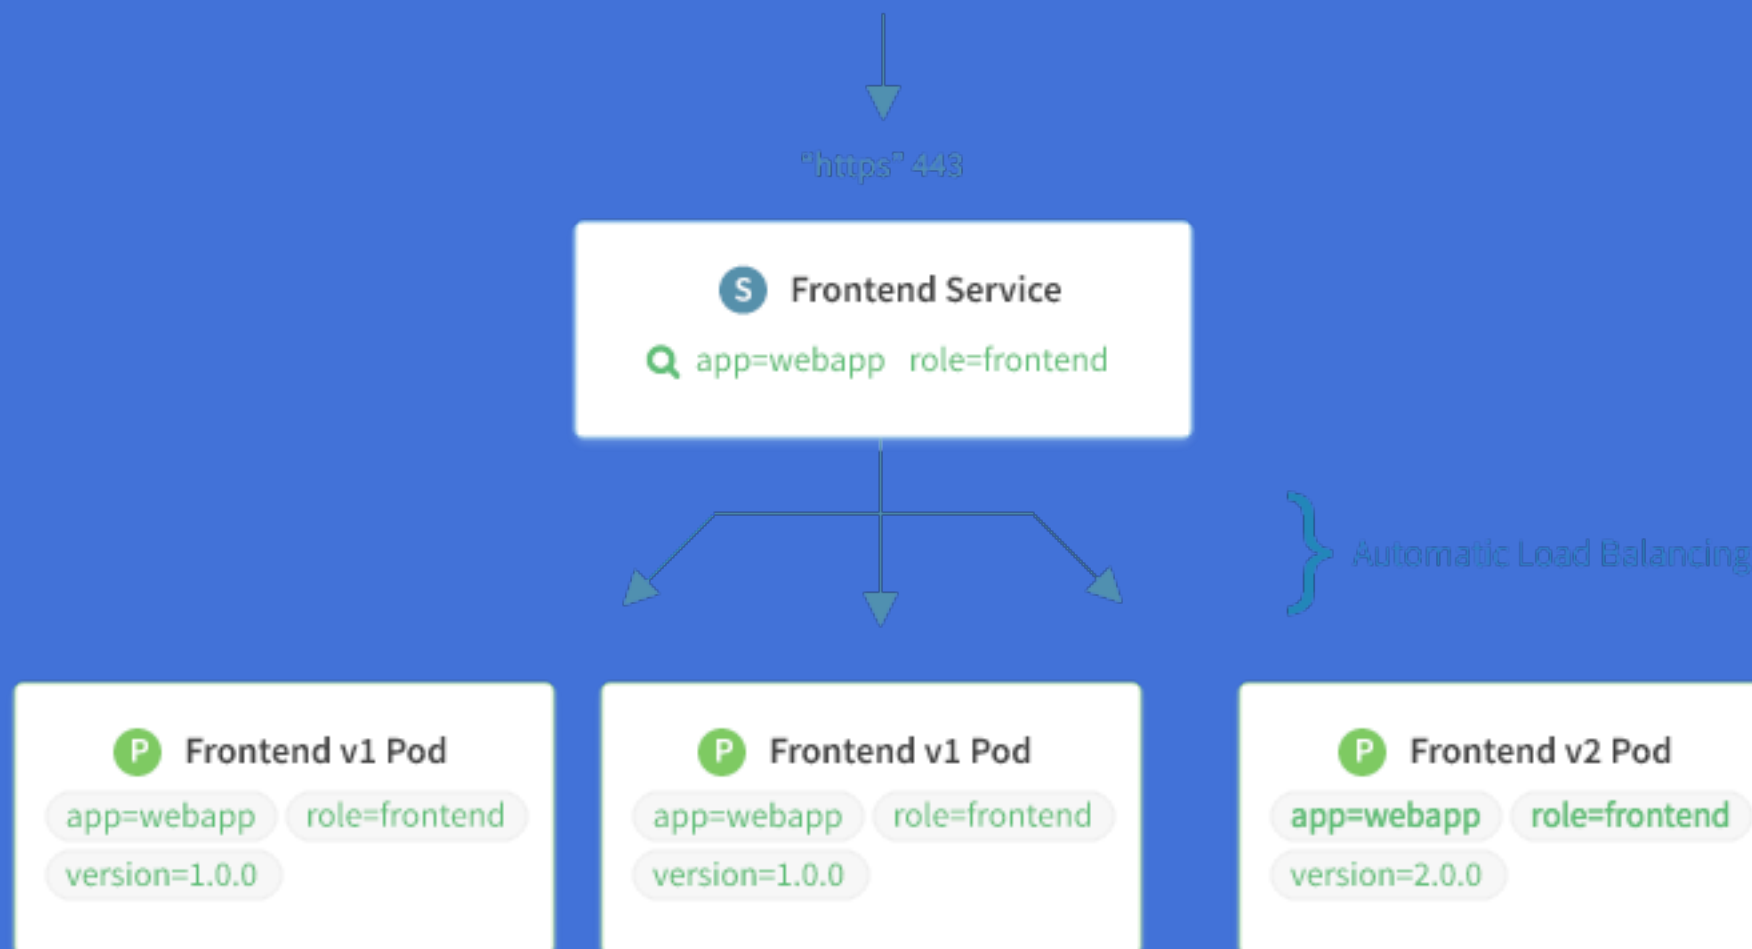
image: tommy351/my-server

ports:

– **containerPort:** 4000

Service

- 把指定的 Pod 暴露出來讓外部存取
- 如果有多個 Pod 的話可進行負載平衡



Service

- 三種服務類型：
 - ClusterIP
 - 預設的服務類型，只有 Cluster 內部能存取
 - NodePort
 - 暴露到某個 port
 - LoadBalancer
 - 暴露到外部 IP（需要雲端服務配合）

Service

apiVersion: v1

kind: Service

metadata:

name: my-server

spec:

type: NodePort

ports:

– **port:** 80

targetPort: 4000

4000 (pod) → 80 (service)

selector:

name: my-server

透過 Pod labels 選擇指定的 Pod

服務探索

- Kubernetes 提供了兩種方式讓 Cluster 內部的 Pod 可以互相找到對方
- DNS
 - Kubernetes 內建了 kube-dns，可以透過 Service name 存取
- 環境變數

服務探索

例：有一個名為 redis 的服務

DNS

redis:6379

環境變數

REDIS_SERVICE_HOST=10.0.0.11

REDIS_SERVICE_PORT=6379

REDIS_PORT=tcp://10.0.0.11:6379

REDIS_PORT_6379_TCP=tcp://10.0.0.11:6379

REDIS_PORT_6379_TCP_PROTO=tcp

REDIS_PORT_6379_TCP_PORT=6379

REDIS_PORT_6379_TCP_ADDR=10.0.0.11

5

Pod Management

資料持久化

- Container 內的資料必須存在外部，否則 Container 終止後就會被刪除
- emptyDir
 - Pod 啟動時分配空間，終止後即被刪除
- hostPath
 - 掛載主機的路徑

資料持久化

- gcePersistentDisk, awsElasticBlockStore, AzureFileVolume, AzureDiskVolume
 - GCE, AWS, Azure 提供的磁碟
- nfs
- persistentVolumeClaim
- <http://kubernetes.io/docs/user-guide/volumes/>

資料持久化

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
  labels:
    name: redis
spec:
  containers:
    - name: redis
      image: redis:3.2
      ports:
        - containerPort: 6379
      volumeMounts:
        - mountPath: /data
          name: data
  volumes:
    - name: data
      emptyDir: {}
```

設定管理

- 環境變數
- ConfigMap
- Secret
- 利用 ConfigMap 和 Secret 的方式：
 - 掛載磁碟
 - 環境變數

環境變數

apiVersion: v1

kind: Pod

metadata:

name: my-server

labels:

name: my-server

spec:

containers:

 - **name:** my-server

image: tommy351/my-server

ports:

 - **containerPort:** 4000

env:

 - **name:** REDIS_ADDRESS

value: redis:6379

 - **name:** SERVER_HOST

value: :4000

REDIS_ADDRESS="redis:6379"

SERVER_HOST=":4000"

ConfigMap

apiVersion: v1

kind: ConfigMap

metadata:

name: redis

data:

redis.conf: |

 maxmemory 2mb

 maxmemorypolicy allkeyslru

Secret

apiVersion: v1

kind: Secret

metadata:

name: redis

type: Opaque

data:

username: YWRtaW4=

password: cGFzcw==

data 必須經過 base64 編碼
password=base64("pass")
username=base64("admin")

掛載 ConfigMap / Secret

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
  labels:
    name: redis
spec:
  containers:
    - name: redis
      image: redis:3.2
      ports:
        - containerPort: 6379
      volumeMounts:
        - mountPath: /usr/local/etc/redis
          name: config
        - mountPath: /srv/redis/secret
          name: secret
  volumes:
    - name: config
      configMap:
        name: redis
    - name: secret
      secret:
        secretName: redis
```

/usr/local/etc/redis/redis.conf

maxmemory 2mb

maxmemorypolicy allkeyslru

/srv/redis/secret/username

admin

/srv/redis/secret/password

pass

從 ConfigMap / Secret 設定環境變數

```
apiVersion: v1
kind: Pod
metadata:
  name: redis
  labels:
    name: redis
spec:
  containers:
  - name: redis
    image: redis:3.2
    ports:
      - containerPort: 6379
    env:
      - name: REDIS_CONF
        valueFrom:
          configMapKeyRef:
            name: redis
            key: redisconf
      - name: REDIS_USERNAME
        valueFrom:
          secretKeyRef:
            name: redis
            key: username
      - name: REDIS_PASSWORD
        valueFrom:
          secretKeyRef:
            name: redis
            key: password
```

REDIS_CONF

maxmemory 2mb

maxmemorypolicy allkeys-lru

REDIS_USERNAME

admin

REDIS_PASSWORD

pass

健康檢查

- 檢查 Pod 是否正常執行
- 兩種檢查時期：
 - livenessProbe
 - 檢查 Pod 是否還活著，每隔一段時間就會執行
 - readinessProbe
 - 確保 Pod 已經開啟並能正常運作，在 Pod 進入 Running 狀態前執行

健康檢查

- 三種檢查方式：
 - exec
 - 執行指令並檢查回傳值是否為 0
 - httpGet
 - 發送 HTTP 請求，並檢查回傳狀態為 2xx
 - tcpSocket
 - 確保 TCP socket 開啟

健康檢查

```
apiVersion: v1
kind: Pod
metadata:
  name: my-server
  labels:
    name: my-server
spec:
  containers:
    - name: my-server
      image: tommy351/my-server
      ports:
        - containerPort: 4000
      livenessProbe:
        initialDelaySeconds: 15
        periodSeconds: 10
        httpGet:
          path: /
          port: 4000
        timeoutSeconds: 5
```

Pod 啟動 15 秒後

每隔 10 秒

戳 <http://localhost:4000/>

如果回傳狀態不是 2xx

或超過 5 秒沒有回應

即判斷 Pod 不健康

Rolling Update

- 從舊的 image 漸進更新到新的 image
 - 建立新的 Replication controller
 - 用新的 image 開啟 container 並漸漸消滅舊的 container
 - 完成後重新命名新的 Replication controller 並消滅舊的

Rolling Update

```
$ kubectl rolling-update my-server --image=tommy351/my-server:v2
```

```
Created my-server-bc33bf3389402cd633ec4573695db4dc
```

```
Scaling up my-server-bc33bf3389402cd633ec4573695db4dc from 0 to 2, scaling down my-server from 2 to 0 (keep 2 pods available, don't exceed 3 pods)
```

```
Scaling my-server-bc33bf3389402cd633ec4573695db4dc up to 1
```

```
Scaling my-server down to 1
```

```
Scaling my-server-bc33bf3389402cd633ec4573695db4dc up to 2
```

```
Scaling my-server down to 0
```

```
Update succeeded. Deleting old controller: my-server
```

```
Renaming my-server-bc33bf3389402cd633ec4573695db4dc to my-server  
replicationcontroller "my-server" rolling updated
```

```
$ kubectl get rc -o wide
```

NAME	DESIRED	CURRENT	AGE	CONTAINER(S)
------	---------	---------	-----	--------------

IMAGE(S)	SELECTOR
----------	----------

my-server	2	2	23s	my-server
-----------	---	---	-----	-----------

```
tommy351/my-server:v2
```

```
deployment=bc33bf3389402cd633ec4573695db4dc,name=my-server
```

Rollback

Rolling update 途中有可能因為 crash 或中斷而失敗，可利用 `--rollback` 回退到上一個版本

```
$ kubectl rolling-update my-server --rollback
```

資源管理

- 可設定 CPU、記憶體限制
- requests
 - 必須有足夠資源才能啟動 Pod，否則會在 Pending 狀態等待
- limits
 - 超過資源上限的話，會使 Pod 終止
 - 如果沒有設定 requests 的話，requests = limit

資源管理

apiVersion: v1

kind: Pod

metadata:

name: my-server

labels:

name: my-server

spec:

containers:

 - **name:** my-server

image: tommy351/my-server

ports:

 - **containerPort:** 4000

resources:

requests:

cpu: 200m

memory: 100Mi

limits:

cpu: 500m

memory: 200Mi

requests.cpu = 0.2 core

requests.memory = 100MB

limits.cpu = 0.5 core

limits.memory = 200MB

水平擴展

- 手動擴展
- 自動擴展
 - Pod 必須設定 CPU requests
 - 根據 CPU 使用量

手動擴展

```
$ kubectl scale rc my-server --replicas=6
```

```
replicationcontroller "my-server" scaled
```

```
$ kubectl get rc
```

NAME	DESIRED	CURRENT	AGE
my-server	6	6	1m

```
$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
my-server-54m5f	1/1	Running	0	26s
my-server-9qxko	1/1	Running	0	26s
my-server-e9bch	1/1	Running	0	26s
my-server-tzjlm	1/1	Running	0	26s
my-server-yo6j8	1/1	Running	0	1m
my-server-zeh6e	1/1	Running	0	1m

自動擴展

```
$ kubectl autoscale rc my-server --cpu-percent=50 --  
min=1 --max=10
```

```
replicationcontroller "my-server" autoscaled
```

HorizontalPodAutoscaler

```
apiVersion: extensions/v1beta1
kind: HorizontalPodAutoscaler
metadata:
  name: my-server
spec:
  scaleRef:
    kind: ReplicationController
    name: my-server
    subresource: scale
minReplicas: 1
maxReplicas: 10
cpuUtilization:
  targetPercentage: 50
```

A

Appendix

Google Cloud

Google Container Engine

- 如果你懶得自己架 Kubernetes cluster 的話，Google Container Engine 大概是最輕鬆簡單的 solution
- Load balancer
- Cloud Logging
- Private Docker Registry
- Cloud Monitoring

Ingress

- 自動在 Google Cloud 上建立 Load balancer
- 必須使用 NodePort service
- Google Cloud 的 Load balancer 也有健康檢查，但是和 Kubernetes 不互通，要另外設定
- 必須暴露 30000-32767 TCP port（NodePort 的範圍）才能讓 Google Cloud 健康檢查

Ingress

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: test
spec:
  rules:
  - host: foo.bar.com
    http:
      paths:
      - path: /foo
        backend:
          serviceName: s1
          servicePort: 80
      - path: /bar
        backend:
          serviceName: s2
          servicePort: 80
```

Cloud Logging

- 不用設定
- 所有 Pod 的紀錄都會自動送到 Google Cloud

按標籤或搜尋字詞挑選

■ 建立指標

Container Engine, staging, default

所有紀錄

任何紀錄層級 ▾

指定日期

►

C

2016-10-01

檢視選項

```

> !! 14:07:00.000 ERROR: syntax error at or near "member_emails"
> !! 14:07:00.000 LINE 1: DELETE member_emails WHERE activated = FALSE and created_at ...
> !! 14:07:00.000 ^
> !! 14:07:01.000 Sat Oct 01 2016 06:07:01 GMT+0000 Thu Aug 04 2016 15:18:00 GMT+0800 5006941
> !! 14:07:01.000 Sat Oct 01 2016 06:07:01 GMT+0000 Thu Aug 04 2016 18:00:00 GMT+0800 4997221
> !! 14:07:01.000 Sat Oct 01 2016 06:07:01 GMT+0000 Mon Aug 08 2016 22:00:00 GMT+0800 4637221
> !! 14:07:01.000 Sat Oct 01 2016 06:07:01 GMT+0000 Tue Aug 09 2016 18:00:00 GMT+0800 4555221
> !! 14:07:01.000 REFRESH MATERIALIZED VIEW
> !! 14:07:01.000 Refreshing latest_posts done.
> !! 14:07:06.000 REFRESH MATERIALIZED VIEW
> !! 14:07:06.000 Refreshing members_full done.
> !! 14:07:10.000 2016/10/01 06:07:10 Output [influxdb] buffer fullness: 28 / 10000 metrics. Total gathered metrics: ...
> !! 14:07:10.000 2016/10/01 06:07:10 Output [influxdb] wrote batch of 28 metrics in 6.87926ms
> !! 14:07:20.000 2016/10/01 06:07:20 Output [influxdb] buffer fullness: 28 / 10000 metrics. Total gathered metrics: ...
> !! 14:07:20.000 2016/10/01 06:07:20 Output [influxdb] wrote batch of 28 metrics in 5.992194ms
> !! 14:07:30.000 2016/10/01 06:07:30 Output [influxdb] buffer fullness: 28 / 10000 metrics. Total gathered metrics: ...
> !! 14:07:30.000 2016/10/01 06:07:30 Output [influxdb] wrote batch of 28 metrics in 6.952921ms
> !! 14:07:40.000 2016/10/01 06:07:40 Output [influxdb] buffer fullness: 28 / 10000 metrics. Total gathered metrics: ...
> !! 14:07:40.000 2016/10/01 06:07:40 Output [influxdb] wrote batch of 28 metrics in 7.410386ms
> !! 14:07:50.000 2016/10/01 06:07:50 Output [influxdb] buffer fullness: 28 / 10000 metrics. Total gathered metrics: ...
> !! 14:07:50.000 2016/10/01 06:07:50 Output [influxdb] wrote batch of 28 metrics in 7.950795ms

```


在符合目前篩選條件的項目中，找不到任何更新的紀錄檔。

1

Thanks!